

Cloud Computing Specialization Capstone – Task 1

Overview

The capstone project involves using cloud computing services to process and analyse a publicly available dataset. The dataset is from the Bureau of Transportation Statistics and we will be using data pertaining to airline traffic to answer a series of questions. In this first task we will use *batch processing* techniques to process the data. Specifically, we will use tools that are included in Apache Hadoop toolset and these will be run on the Amazon Web Services (AWS) platform.

A key consideration throughout this project will be the efficient and appropriate use of the various services that are offered by AWS. The services must be combined in a way that maximizes ease of use and minimizes costs. AWS services often have a 'Free Tier', that allow you to use that service up to some limit e.g. storage, run-time etc. and we aim to minimize the costs that are incurred as a result of exceeding these limits.

Cleaning and storing the data

The transportation dataset was initially stored as a 15GB Elastic Block Storage (EBS) snapshot. The snapshot was used to create an Elastic Compute (EC2) volume and was then mounted to an EC2 instance. The 'aviation' subfolder was uploaded to Amazon's cheap, long term storage solution called S3. Amazon S3 offers 5GB of free storage and by uploading only the 3.9GB aviation arrival/departure time data, we remained below this threshold. The data was then downloaded to a local machine, where it could be explored and cleaned before being re-uploaded to S3.

We used a shell script to unzip the many downloaded files on the local machine. The Pandas package for Python was used to clean the files sequentially and produce a cleaned set of files. Records were dropped if the flight was cancelled or they were an exact duplicate of another record. There was a significantly amount of rows dropped due to duplication. Subsequently, all columns that were not necessary for answering the questions for task 1 were dropped. Then, any rows that had any missing values in the remaining columns were dropped. The cleaning procedure reduced the file sizes by a factor of ~7. The files were then all aggregated into one file by appending 30 files worth of data at a time in order to avoid excessive memory usage. The total uncompressed file size was ~5.8GB. The aggregated and cleaned .csv file was uploaded back to S3 storage.

AWS Architecture

Amazon EMR was configured to run with the default three node configuration using m5.xlarge instance types - 4 vCore, 16 GB memory, 64 GB EBS storage. These instances were estimated to have sufficient memory to hold our PySpark dataframes based on testing on a local machine with 16GB memory. One instance operated as the Master node type and the remaining two nodes operated as Core node types. We used EMR release version 5.36.0 and had the following services added to our configuration: Hadoop 2.10.1, JupyterHub 1.4.1, Spark 2.4.8 and Livy 0.7.1. We used a bootstrap .sh file to install the necessary Python 3 modules on our instances.

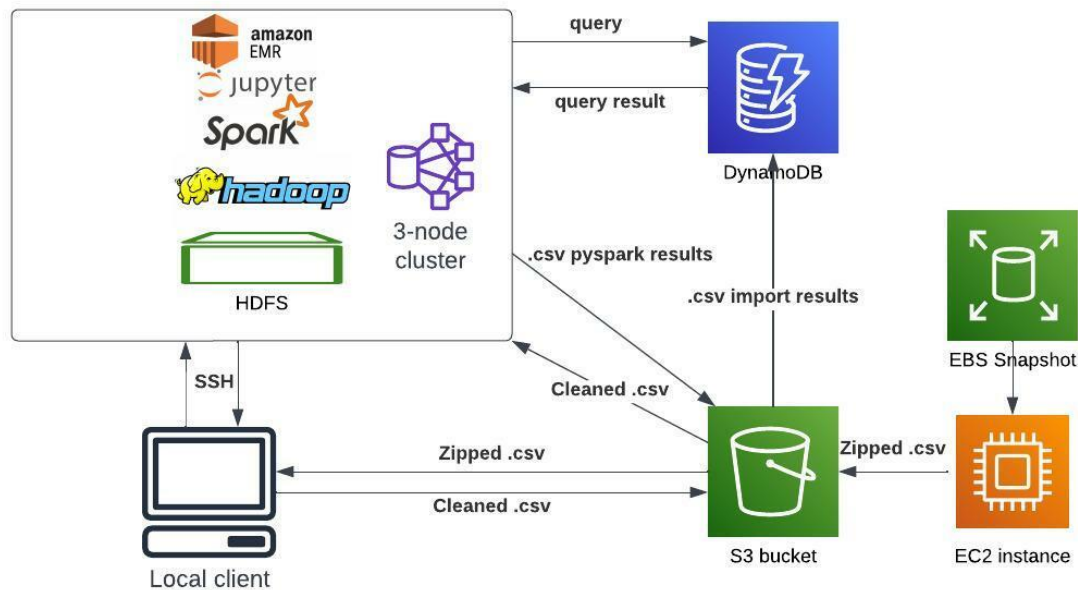


Figure 1. AWS architecture with data flow

General Data Analysis Strategy

We decided to use PySpark within a Jupyter Notebook to do our data analysis, due to familiarity with the Python programming language. The cleaned data was uploaded to the EMR cluster from S3 and was loaded into a PySpark dataframe. Once the data was in a dataframe we could perform necessary data manipulation to answer the questions for Task 1. Before launching the cluster, the Python code was tested locally on a smaller data subset to ensure correctness and to informally observe the performance of different methods. The strategy helped to reduce the amount of testing required once the cluster was running, which had the result of reducing the total running time of the cluster so that costs could be minimized. It was found that using the `.show()` method on dataframes was slow, so was used whilst debugging and otherwise only when necessary to produce the final result. Similarly, when producing `.csv` files of the results for the necessary questions, using the `.coalesce(1)` and `.toPandas()` functions required higher processing time and could not be used on Question 3.2 due to the size of the dataframe. Instead, on this last part we used a separate Python script to combine the `.csv` files produced by the `.to_csv()` method, which was still slow. Producing these `.csv` files was necessary to import the data into DynamoDB due to difficulty with transferring the data directly from the PySpark dataframes. In this way, using Hive would have been simpler, but would have required more time to gain familiarity with.

We found that some of the results differed slightly from the published reference result, although generally not by much. The difference is likely due to the presence of several corrupted `.csv` files from the 2008 series and due to a difference in cleaning techniques. In particular, we removed many duplicate flight records from the initial data set and it could be that this was not done with the reference results.

Group 1 Question 1

```
df_pyspark.groupBy('Origin').count().orderBy('count', ascending = [False]).show(n=10)
```

Group 1 Question 2

```
df_pyspark.groupBy('Carrier').avg('ArrDelay').orderBy(['avg(ArrDelay)'], ascending = [True]).show(n=10)
```

Group 1 Question 3

```
df_pyspark.groupBy('DayOfWeek').avg('ArrDelay').orderBy(['avg(ArrDelay)'], ascending = [True]).show(n=10)
```

G1Q1		G1Q3		G1Q2	
Origin	Count	DayOfWeek	avg(ArrDelay)	Carrier	avg(ArrDelay)
ORD	5998002	Saturday	4.30165	HA	-1.0118
ATL	5659166	Tuesday	5.990445	AQ	1.156923
DFW	5270525	Sunday	6.61328	PS	1.450639
LAX	3788336	Monday	6.716091	ML	4.747609
PHX	3246442	Wednesday	7.203647	PA	5.322431
DEN	3078009	Thursday	9.094438	F9	5.466049
DTW	2738373	Friday	9.720996	NW	5.557783
IAH	2700746			WN	5.560774
MSP	2533996			OO	5.735951
SFO	2530684			9E	5.867185

Group 2 Question 1

```
window = Window.partitionBy('Origin').orderBy('avg(DepDelay)')
df_dep_perf_by_carrier = df_pyspark.groupBy('Origin', 'Carrier').avg('DepDelay')\
.withColumn('row_number', pyfuncs.row_number().over(window))
df_dep_perf_by_carrier = df_dep_perf_by_carrier.filter(df_dep_perf_by_carrier['row_number'] <= 10).drop('row_number')
```

```
q2_airports = ['CMI', 'BWI', 'MIA', 'LAX', 'IAH', 'SFO']
df_dep_perf_by_carrier.filter(df_dep_perf_by_carrier.Origin.isin(q2_airports)).show(n=100)
```

BWI		CMI		IAH		LAX		MIA		SFO	
Carrier	avg(DepDelay)	Carrier	avg(DepDelay)	Carrier	avg(DepDelay)	Carrier	avg(DepDelay)	Carrier	avg(DepDelay)	Carrier	avg(DepDelay)
F9	0.746493	OH	0.612005	NW	3.515589	RU	1.948387	9E	-3.66667	TZ	3.935139
PA	4.761905	US	1.82441	PA	3.901583	MQ	2.399978	EV	1.202643	MQ	4.782694
CO	5.113806	TW	3.744589	PI	3.952167	OO	4.214037	RU	1.283429	F9	5.155881
YV	5.504666	PI	4.369928	RU	4.778435	FL	4.687786	TZ	1.782244	PA	5.283914
NW	5.655729	DH	6.08016	AL	4.962449	TZ	4.758884	XE	2.745645	NW	5.717548
AL	5.728207	EV	6.665138	US	5.031701	PS	4.805362	PA	4.026508	PS	6.256748
AA	5.913444	MQ	7.992782	F9	5.549463	NW	5.092928	NW	4.423478	DL	6.512475
9E	7.220703			AA	5.616927	F9	5.719679	US	5.990862	CO	7.047357
US	7.471462			TW	6.027508	HA	5.817645	UA	6.808275	US	7.347713
DL	7.641127			WN	6.213109	YV	6.036279	ML	7.512146	TW	7.763071

Group 2 Question 2

```

window = Window.partitionBy('Origin').orderBy('avg(DepDelay)')
df_dep_perf_by_dest = df_pyspark.groupBy('Origin', 'Dest').avg('DepDelay')\
.withColumn('row_number', pyfunctools.row_number().over(window))
df_dep_perf_by_dest = df_dep_perf_by_dest.filter(df_dep_perf_by_dest['row_number'] <= 10).drop('row_number')
q2_airports = ['CMI', 'BWI', 'MIA', 'LAX', 'IAH', 'SFO']
df_dep_perf_by_dest.filter(df_dep_perf_by_dest['Origin'].isin(q2_airports)).show(n=100)

```

BWI		CMI		LAX		IAH		SFO		MIA	
Dest	avg(DepDelay)	Dest	avg(DepDelay)	Dest	avg(DepDelay)	Dest	avg(DepDelay)	Dest	avg(DepDelay)	Dest	avg(DepDelay)
MLB	1.161473	PIT	1.107826	SDF	-16	MSN	-2	SDF	-10	SAN	1.710383
DAB	1.508475	CVG	1.902128	LAX	-2	MLI	-1	MSO	-4	BUF	2
SRQ	1.550026	DAY	2.928049	BZN	-0.72727	AGS	-0.6173	LGA	-1.75758	SLC	2.410788
IAD	1.797203	PIA	3.412568	MAF	0	EFD	1.887708	PIE	-1.34104	HOU	2.888491
UCA	3.325903	STL	3.841727	MFE	1.177515	HOU	2.172037	OAK	-0.81371	ISP	3.649123
CHO	3.744928	DFW	5.978074	IYK	1.269825	JAC	2.422619	FAR	0	MEM	3.686327
BGM	3.795134	ATL	6.665138	MEM	1.867979	MTJ	2.902516	BNA	2.379209	GNV	3.960685
DCA	3.970642	ORD	8.163504	PIE	1.926209	RNO	3.227283	MEM	3.262416	PSE	3.975845
GSP	4.192835			SNA	2.094492	GUC	3.537374	SJC	3.985614	TLH	4.234642
SJU	4.203397			SBA	2.274377	BPT	3.599533	MKE	3.988028	MCI	4.612245

Group 2 Question 4

```

df_g2q4 = df_pyspark.groupBy('Origin', 'Dest').avg('ArrDelay').orderBy(['avg(ArrDelay)'], ascending = [True])
q4_pairs = [('CMI', 'ORD'), ('IND', 'CMH'), ('DFW', 'IAH'), ('LAX', 'SFO'), ('JFK', 'LAX'), ('ATL', 'PHX')]
for pair in q4_pairs:
    print(pair)
    df_g2q4.filter((df_g2q4['Origin'] == pair[0]) & (df_g2q4['Dest'] == pair[1])).show(n=10)

```

Origin	Dest	avg(ArrDelay)
CMI	ORD	10.14366
IND	CMH	2.899904
DFW	IAH	7.654443
LAX	SFO	9.589283
JFK	LAX	6.635119
ATL	PHX	9.021342

Group 3 Question 1

```

airport_freqs = df_pyspark.groupBy('Origin').count().select("count").rdd.flatMap(lambda x: x).collect()

```

```

x = np.sort(airport_freqs)
n = len(airport_freqs)
y = 1 - np.arange(n) / float(n)

def power_func(x, a, b):
    return a * x ** (-b)


popt, _ = curve_fit(power_func, x, y)
power_y = power_func(x, *popt)
plt.scatter(x, y, marker='o', facecolors='none', edgecolors='black')
plt.plot(x, power_y, color='r')
shape, loc, scale = lognorm.fit(x, 2)
lognorm_y = lognorm.sf(x, shape, loc=loc, scale=scale)
plt.plot(x, lognorm_y, color='b')
plt.xlabel('x')
plt.ylabel('CCDF')
plt.yscale('log')
plt.xscale('log')
plt.legend(['Data', 'Power', 'Log-normal'])
print('Data - Power Kolmogorov-Smirnov:')
print(ks_2samp(y, power_y))
print('Data - Log-normal Kolmogorov-Smirnov:')
print(ks_2samp(y, lognorm_y))
print('CCDF of data and distributions:')

```

```

Data - Power Kolmogorov-Smirnov:
KstestResult(statistic=0.28034682080924855, pvalue=2.2333610336795968e-12)
Data - Log-normal Kolmogorov-Smirnov:
KstestResult(statistic=0.06647398843930635, pvalue=0.42964381800204454)

```



We plotted the Complementary Cumulative Distribution Function (CCDF) of the number of flights from each airport using log-log scales. Using the SciPy library we fitted a Power-law and a Log-normal curve to the airport popularity data. From the plot we can clearly see that the CCDF does not follow the Power-law curve, but does seem to follow a Log-normal curve. We confirmed this using a Kolmogorov-Smirnov test which had a P-value of < 0.05 for the Power-law test, but a P-value of 0.43

for the Log-normal curve. Using a cut-off of $P = 0.05$, we reject the null hypothesis for the Power-law curve and conclude that it is significantly different to the data. We fail to reject the null hypothesis for the log-normal curve and conclude that it is not significantly different to the data.

Group 3 Question 2

```
df_pyspark = df_pyspark.withColumn('Year', date_format('FlightDate', 'y'))
df_2008 = df_pyspark.filter(df_pyspark['Year'] == 2008).drop('Year')
df_2008.show(n=10)
w = Window.partitionBy('FlightDate', 'Origin', 'Dest')

df_bestXY = df_2008.alias('df1').filter(df_2008['DepTime'] < 1200)
df_bestXY = df_bestXY.withColumn('minArrDelay', pyfuncs.min('ArrDelay').over(w))\
.where(col('ArrDelay') == col('minArrDelay')).drop('minArrDelay')
df_bestYZ = df_2008.alias('df2').filter(df_2008['DepTime'] >= 1200)
df_bestYZ = df_bestYZ.withColumn('minArrDelay', pyfuncs.min('ArrDelay').over(w))\
.where(col('ArrDelay') == col('minArrDelay')).drop('minArrDelay')
col_names = df_2008.columns
df_bestXY = df_bestXY.selectExpr([col + ' as XY_' + col for col in col_names])
df_bestYZ = df_bestYZ.selectExpr([col + ' as YZ_' + col for col in col_names])
df_g3q2 = df_bestXYZ\
.select('XY_FlightDate', 'XY_Origin', 'XY_Dest', 'XY_ArrDelay', 'YZ_FlightDate', 'YZ_Origin', 'YZ_Dest', 'YZ_ArrDelay')
df_g3q2.write.option('header', 'true').csv("./g3q2")
routes = [
    ('CMI', 'ORD', 'LAX', datetime.datetime(2008, 3, 4, 0, 0, 0)),
    ('JAX', 'DFW', 'CRP', datetime.datetime(2008, 9, 9, 0, 0, 0)),
    ('SLC', 'BFL', 'LAX', datetime.datetime(2008, 4, 1, 0, 0, 0)),
    ('LAX', 'SFO', 'PHX', datetime.datetime(2008, 7, 12, 0, 0, 0)),
    ('DFW', 'ORD', 'DFW', datetime.datetime(2008, 6, 10, 0, 0, 0)),
    ('LAX', 'ORD', 'JFK', datetime.datetime(2008, 1, 1, 0, 0, 0))
]
for route in routes:
    df_bestXYZ.where((df_bestXYZ.XY_FlightDate == route[3]) & (df_bestXYZ.XY_Origin == route[0]) \
    & (df_bestXYZ.XY_Dest == route[1]) & (df_bestXYZ.YZ_Dest == route[2])).show()
```

XY_FlightDate	XY_Carrier	XY_Origin	XY_Dest	XY_DepTime	XY_ArrDelay	YZ_FlightDate	YZ_Carrier	YZ_Dest	YZ_DepTime	YZ_ArrDelay
04/03/2008	MQ	CMI	ORD	710	-14	06/03/2008	AA	LAX	1952	-24
09/09/2008	AA	JAX	DFW	722	1	11/09/2008	MQ	CRP	1648	-7
01/04/2008	OO	SLC	BFL	1101	12	03/04/2008	OO	LAX	1509	6
12/07/2008	WN	LAX	SFO	650	-13	14/07/2008	US	PHX	1916	-19
10/06/2008	UA	DFW	ORD	658	-21	12/06/2008	AA	DFW	1650	-10
01/01/2008	UA	LAX	ORD	700	1	03/01/2008	B6	JFK	1853	-7

Conclusion

Using Jupyter notebooks and PySparks was a convenient and easy to use methodology. By using the local machine to extensively test code before using it on the EMR cluster we significantly reduced running costs. The downside of this methodology was the minor difficulty in getting the results that had to be placed into persistent storage to DynamoDB. Using intermediate .csv files, uploading them to s3 and then importing them to DynamoDB was not difficult, but it was a relatively slow procedure. Other options included using boto3's batch writing, which would require sequentially going through the PySpark dataframes, which would also be slow. Using Hive may have been more efficient in this regard, but unfamiliarity with the framework made this a less attractive option.